

Interpretability, VAEs, and GANs

Recitation 3

MIT - 6.802 / 6.874 / 20.390 / 20.490 / HST.506 - Spring 2021

Dylan Cable

Variational Inference

Slide credit to – Jeff Miller, Harvard Biostatistics

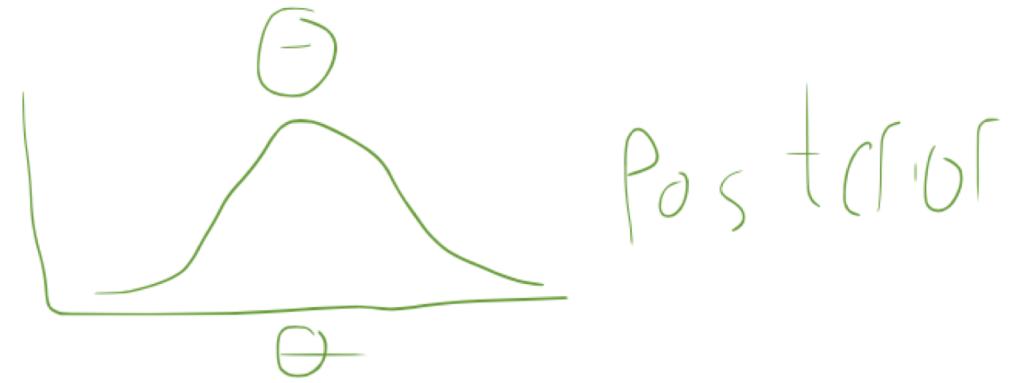
$$p(\theta | x)$$

Params Data

Posterior

$$p(\theta | x) \propto p(\theta)p(x|\theta)$$

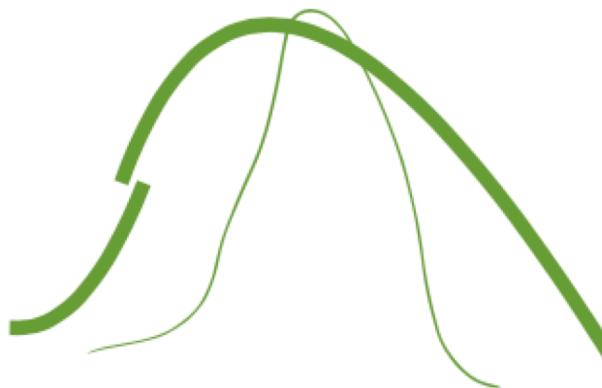
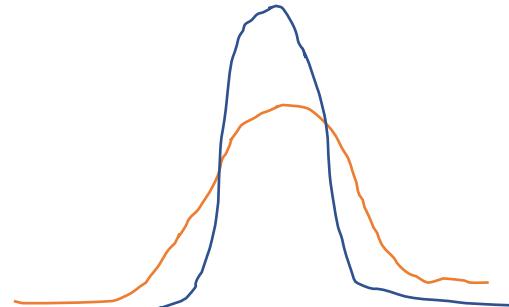
$p(x|\theta)$ Likelihood



Posterior

Introduction

- Variational inference (VI) is an approach to posterior inference based on approximating the posterior by a nicer distribution.
- The general idea is to:
 1. choose a nice family of distributions \mathcal{Q} ,
 2. find a $q \in \mathcal{Q}$ that is as close as possible to the posterior, and
 3. use q to quantify uncertainty, as a proxy for the posterior.
- Different choices of \mathcal{Q} and definitions of “close” lead to different variational inference techniques.
- In Bayesian statistics, VI is also known as variational Bayes, but VI is also useful outside of Bayesian inference.



Likelihood Ratio Test

$$\log\left(\frac{g_f(x)}{\pi(x)}\right) \geq T$$

← optimal test

$$P(g_f || \pi) = \int g_f(x) \log\left(\frac{g_f}{\pi}\right)$$

Classic variational inference

- Let $\pi(\theta)$ denote the target distribution, e.g., $\pi(\theta) = p(\theta|x)$.
- Classic VI makes the following choices:

1. the approximating family \mathcal{Q} consists of all factorized distributions of the form

$$q(\theta) = q_1(\theta_1) \cdots q_m(\theta_m)$$

for some decomposition of θ into components $\theta_1, \dots, \theta_m$, and

2. we seek the $q \in \mathcal{Q}$ that minimizes the Kullback–Leibler divergence from q to π ,

$$q^{\text{opt}} \in \underset{q \in \mathcal{Q}}{\operatorname{argmin}} D(q\|\pi).$$

- The *Kullback–Leibler (KL) divergence*, or *relative entropy*, is

$$D(q\|\pi) = \int q(\theta) \log \frac{q(\theta)}{\pi(\theta)} d\theta = \mathbb{E}_q \left(\log \frac{q(\theta)}{\pi(\theta)} \right).$$

Toy example: VI for univariate normal

- Consider the univariate normal model as a toy example:

$$X_1, \dots, X_n | \mu, \lambda \stackrel{\text{iid}}{\sim} \mathcal{N}(\mu, \lambda^{-1})$$

and assume an improper uniform prior on μ and λ .

- Define the target distribution to be the posterior:

$$\begin{aligned}\pi(\mu, \lambda) &= p(\mu, \lambda | x_{1:n}) \propto p(x_{1:n} | \mu, \lambda) \\ &\propto_{\mu, \lambda} \lambda^{n/2} \exp\left(-\frac{1}{2}\lambda \sum_{i=1}^n (x_i - \mu)^2\right).\end{aligned}$$

- Thus, $\log \pi(\mu, \lambda) = \frac{n}{2} \log \lambda - \frac{1}{2}\lambda \sum_{i=1}^n (x_i - \mu)^2 + \text{const.}$
- A natural decomposition to try would be $\theta_1 = \mu$ and $\theta_2 = \lambda$, that is, to consider approximations of the form

$$q(\mu, \lambda) = q_1(\mu)q_2(\lambda).$$

For convenience, we write $q(\mu, \lambda) = q(\mu)q(\lambda)$.

Classic variational inference: Algorithm

- To derive the classic VI updates, we need only specify the target distribution π and the decomposition of θ into $\theta_1, \dots, \theta_m$.
- The classic VI algorithm then proceeds as follows:
 1. Initialize q_1, \dots, q_m .
 2. Repeat until convergence:
For $j = 1, \dots, m$, update q_j by setting it to be

$$q_j^{\text{new}}(\theta_j) \propto \exp(h_j(\theta_j))$$

where

$$h_j(\theta_j) := \mathbb{E}_q(\log \pi(\theta) \mid \theta_j) = \int (\log \pi(\theta)) \prod_{i \neq j} q_i(\theta_i) d\theta_i.$$

3. Use $q(\theta) = q_1(\theta_1) \cdots q_m(\theta_m)$ as an approximation to $\pi(\theta)$.

Justification of classic VI algorithm (1/2)

- Let $\theta = (\theta_1, \dots, \theta_m)$ where $\theta_i \sim q_i$ independently for some q_1, \dots, q_m . Let π be the target distribution. Then for any j ,

$$\begin{aligned} D(q_1 \cdots q_m \| \pi) &= \mathbb{E} \left(\log \frac{q_1(\theta_1) \cdots q_m(\theta_m)}{\pi(\theta)} \right) \\ &= \sum_{i=1}^m \mathbb{E} \log q_i(\theta_i) - \mathbb{E} \log \pi(\theta) \\ &= \mathbb{E} \log q_j(\theta_j) - \mathbb{E}(\mathbb{E}(\log \pi(\theta) | \theta_j)) + (\text{const wrt } q_j) \\ &= \mathbb{E} \log q_j(\theta_j) - \mathbb{E} h_j(\theta_j) + (\text{const wrt } q_j) \\ &= \mathbb{E} \left(\log \frac{q_j(\theta_j)}{c e^{h_j(\theta_j)}} \right) + (\text{const wrt } q_j) \\ &= D(q_j \| c e^{h_j}) + (\text{const wrt } q_j) \end{aligned}$$

where $h_j(\theta_j) := \mathbb{E}(\log \pi(\theta) | \theta_j)$ and $1/c = \int e^{h_j(\theta_j)} d\theta_j$.

Justification of classic VI algorithm (2/2)

- Hence,

$$\operatorname{argmin}_{q_j} D(q_1 \cdots q_m \| \pi) = \operatorname{argmin}_{q_j} D(q_j \| c e^{h_j}).$$

- By the properties of KL divergence, $D(q_j \| c e^{h_j}) \geq 0$ with equality if and only if $q_j = c e^{h_j}$ almost everywhere.
- Therefore, choosing $q_j \propto e^{h_j}$ minimizes $D(q_1 \cdots q_m \| \pi)$ with respect to q_j , given q_i for $i \neq j$.
- This shows that the classic VI algorithm performs coordinate descent, updating q_j to minimize KL at each step, holding q_i fixed for $i \neq j$.

Classic variational inference: Algorithm

- To derive the classic VI updates, we need only specify the target distribution π and the decomposition of θ into $\theta_1, \dots, \theta_m$.
- The classic VI algorithm then proceeds as follows:
 1. Initialize q_1, \dots, q_m .
 2. Repeat until convergence:
For $j = 1, \dots, m$, update q_j by setting it to be

$$q_j^{\text{new}}(\theta_j) \propto \exp(h_j(\theta_j))$$

where

$$h_j(\theta_j) := \mathbb{E}_q(\log \pi(\theta) \mid \theta_j) = \int (\log \pi(\theta)) \prod_{i \neq j} q_i(\theta_i) d\theta_i.$$

3. Use $q(\theta) = q_1(\theta_1) \cdots q_m(\theta_m)$ as an approximation to $\pi(\theta)$.

Toy example: VI for univariate normal

- Consider the univariate normal model as a toy example:

$$X_1, \dots, X_n | \mu, \lambda \stackrel{\text{iid}}{\sim} \mathcal{N}(\mu, \lambda^{-1})$$

and assume an improper uniform prior on μ and λ .

- Define the target distribution to be the posterior:

$$\begin{aligned}\pi(\mu, \lambda) &= p(\mu, \lambda | x_{1:n}) \propto p(x_{1:n} | \mu, \lambda) \\ &\propto_{\mu, \lambda} \lambda^{n/2} \exp\left(-\frac{1}{2}\lambda \sum_{i=1}^n (x_i - \mu)^2\right).\end{aligned}$$

- Thus, $\log \pi(\mu, \lambda) = \frac{n}{2} \log \lambda - \frac{1}{2}\lambda \sum_{i=1}^n (x_i - \mu)^2 + \text{const.}$
- A natural decomposition to try would be $\theta_1 = \mu$ and $\theta_2 = \lambda$, that is, to consider approximations of the form

$$q(\mu, \lambda) = q_1(\mu)q_2(\lambda).$$

For convenience, we write $q(\mu, \lambda) = q(\mu)q(\lambda)$.

Normal example: Deriving the VI updates (1/3)

- Updating $q(\mu)$ given $q(\lambda)$:

$$\begin{aligned} h_1(\mu) &= \int q(\lambda) \log \pi(\mu, \lambda) d\lambda \\ &= \frac{n}{2} \int q(\lambda) \log \lambda - \frac{1}{2} \left(\sum_i (x_i - \mu)^2 \right) \int \lambda q(\lambda) d\lambda + \text{const.} \end{aligned}$$

- Therefore, according to the algorithm, we update $q(\mu)$ to be

$$\begin{aligned} q^{\text{new}}(\mu) &\propto \exp(h_1(\mu)) \propto \exp \left(-\frac{1}{2} E(\lambda) \sum_i (x_i - \mu)^2 \right) \\ &\underset{\mu}{\propto} \mathcal{N}(\mu \mid \bar{x}, (nE(\lambda))^{-1}). \end{aligned}$$

- Computationally, we only need to compute and store \bar{x} and $E(\lambda)$.
- Here, $E(\lambda) = \int \lambda q(\lambda) d\lambda$ is computed using the current $q(\lambda)$.

Normal example: Deriving the VI updates (2/3)

- Updating $q(\lambda)$ given $q(\mu)$:

$$\begin{aligned} h_2(\lambda) &= \int q(\mu) \log \pi(\mu, \lambda) d\mu \\ &= \frac{n}{2} \log \lambda - \frac{1}{2} \lambda \sum_i \int (x_i - \mu)^2 q(\mu) d\mu + \text{const.} \end{aligned}$$

- Therefore, according to the algorithm, we update $q(\lambda)$ to be

$$\begin{aligned} q^{\text{new}}(\lambda) &\propto \exp(h_2(\lambda)) \propto \lambda^{n/2} \exp(-\frac{1}{2}S\lambda) \\ &\propto_{\lambda} \text{Gamma}(\lambda \mid n/2 + 1, S/2) \end{aligned}$$

where, after plugging in $q(\mu) = \mathcal{N}(\mu \mid \bar{x}, (nE(\lambda))^{-1})$ and simplifying,

$$S = \sum_i \int (x_i - \mu)^2 q(\mu) d\mu = n\hat{\sigma}^2 + 1/E(\lambda).$$

- Here, $\hat{\sigma}^2 = \frac{1}{n} \sum_i (x_i - \bar{x})^2$.

Normal example: Deriving the VI updates (3/3)

- Thus, the updates to $q(\mu)$ and $q(\lambda)$ are:

$$q^{\text{new}}(\mu) = \mathcal{N}(\mu \mid \bar{x}, (nE(\lambda))^{-1}),$$

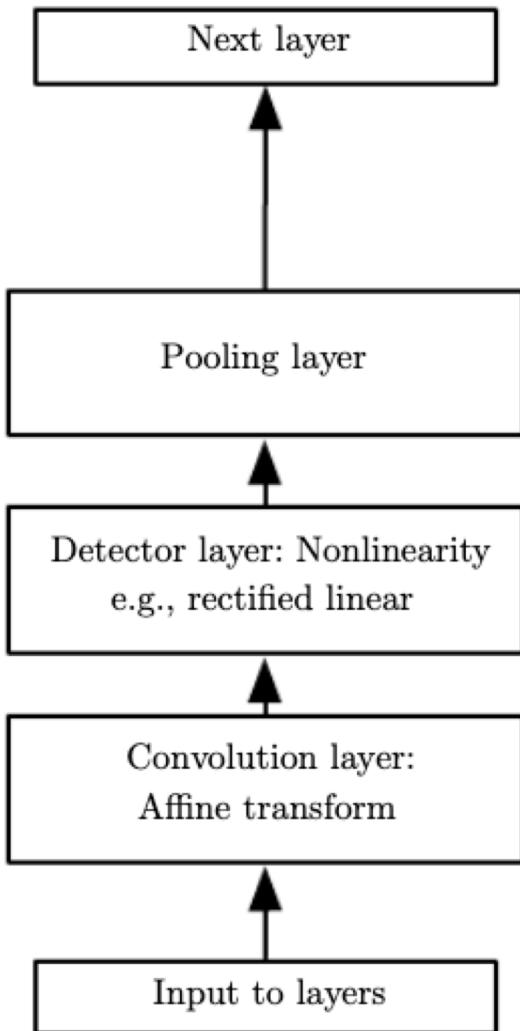
$$q^{\text{new}}(\lambda) = \text{Gamma}(\lambda \mid n/2 + 1, \frac{1}{2}(n\hat{\sigma}^2 + 1/E(\lambda))).$$

- The only thing we need to compute at each iteration is $E(\lambda)$.
- From the form of $q^{\text{new}}(\lambda)$, we see that

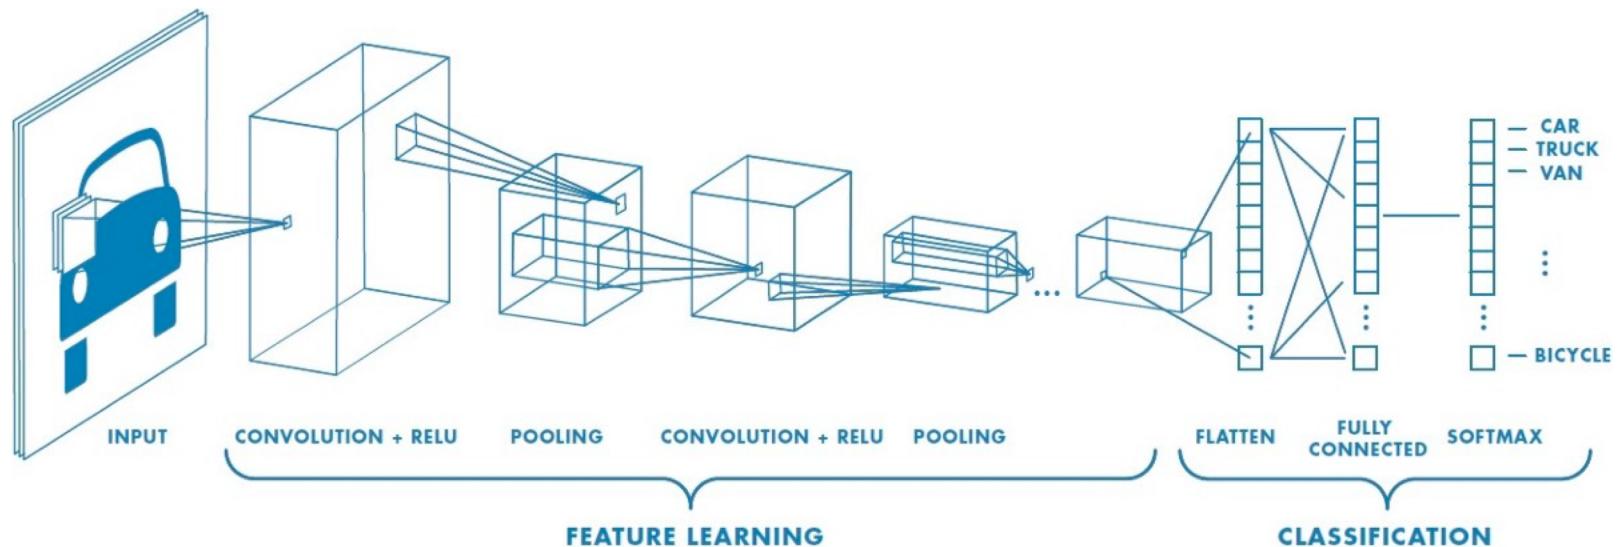
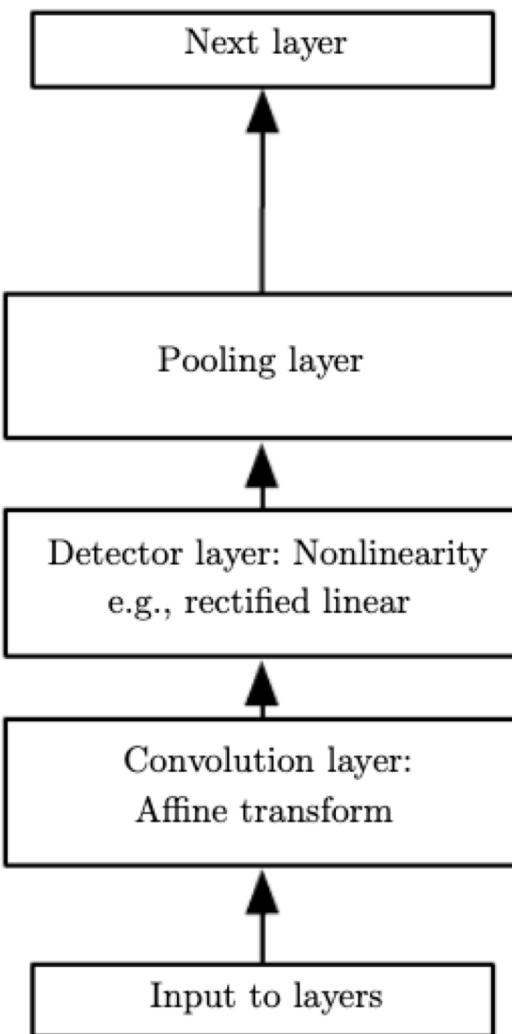
$$E^{\text{new}}(\lambda) = \frac{n/2 + 1}{\frac{1}{2}(n\hat{\sigma}^2 + 1/E(\lambda))}.$$

- In this example, it turns out that we can analytically solve for the limiting value of $E(\lambda)$, which is $E(\lambda) = (n + 1)/(n\hat{\sigma}^2)$.
- However, in more complex models, we will have to iterate until convergence.

Convolutional neural networks: overview



Convolutional neural networks: overview



Convolutional neural networks: convolution operation

Imagine a position x taken at a time point t , which are both continuous,

$$x(t)$$

Suppose this measurement is noisy. To obtain a less noisy estimate of the position, we want a weighted average of recent measurements. Using a weighting function $w(a)$, where a is the age of the measurement,

$$s(t) = \int x(a)w(t-a)da = (x * w)(t)$$

input
convolutional kernel/filter

Generally, sensor inputs are not continuous and most signal processing tasks rely on discretized data—data provided at regular intervals. In machine learning, input is generally a multi-dimensional array of data and the kernel is a multi-dimensional array of parameters:

$$S(i, j) = (W * I)(i, j) = \sum_m \sum_n I(i + m, j + n)W(m, n)$$

This operation also exhibits translational equivariance. Let g be a function mapping one image function to another image function that shifts inputs to the right by one,

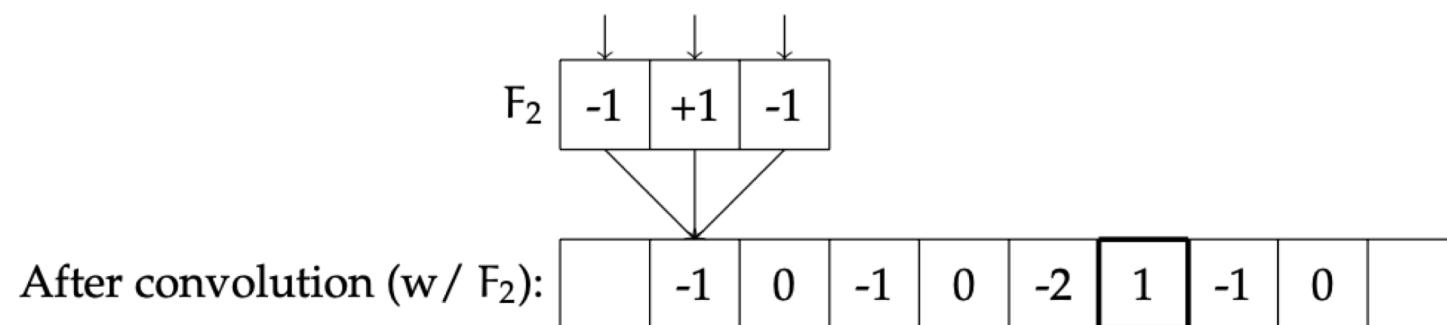
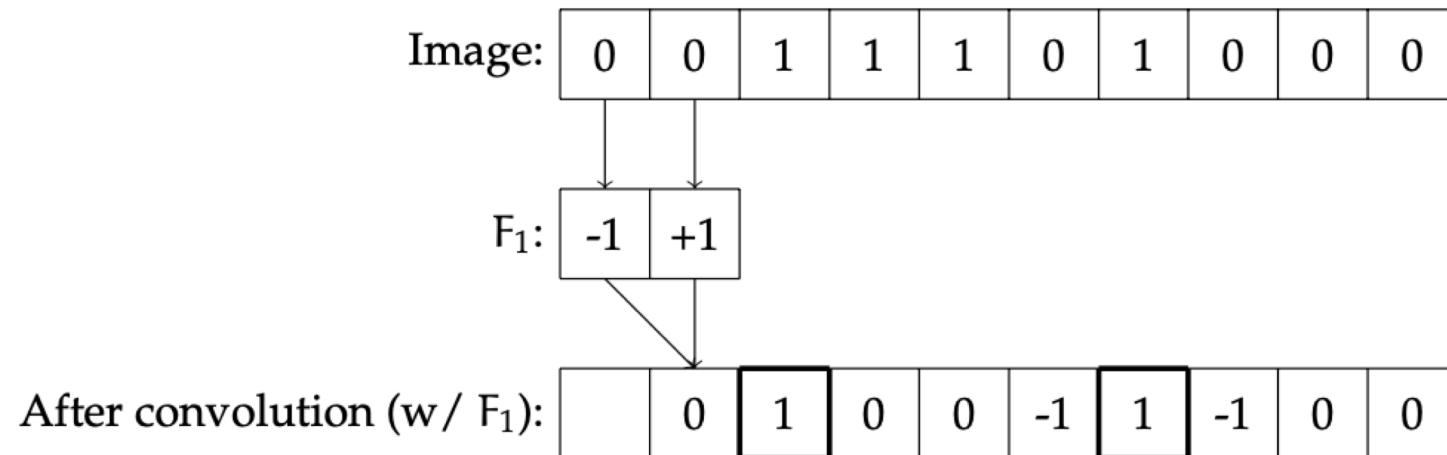
$$I'(i, j) = g(I(i, j)) = I(i - 1, j)$$

If we apply the transformation to I , then apply a convolution, the output is equivalent to applying a convolution to I' then applied the transformation g to the output,

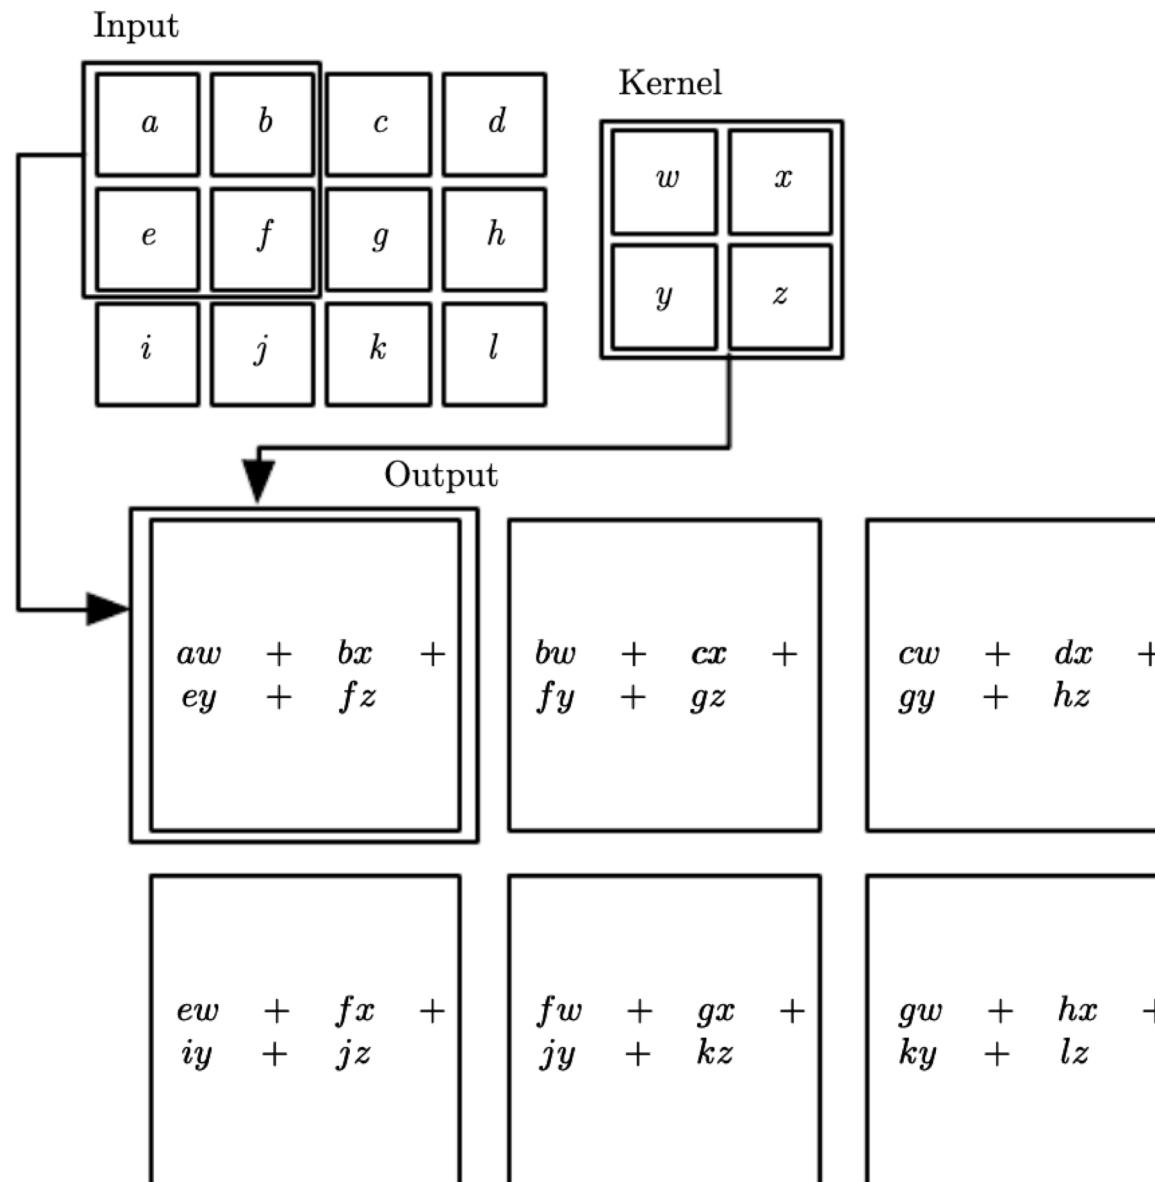
$$(W * I')(i, j) = g((W * I)(i, j))$$

So, a filter applied over an input pattern that is shifted in an image will result in the same convolutional output as the same translational transformation applied after the convolutional—translational equivariance.

Convolutional neural networks: 1D convolution



Convolutional neural networks: 2D convolution



Convolutional neural networks: 2D convolution

1	0	1
0	1	0
1	0	1

Filter / Kernel

1 <small>$\times 1$</small>	1 <small>$\times 0$</small>	1 <small>$\times 1$</small>	0	0
0 <small>$\times 0$</small>	1 <small>$\times 1$</small>	1 <small>$\times 0$</small>	1	0
0 <small>$\times 1$</small>	0 <small>$\times 0$</small>	1 <small>$\times 1$</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

Convolutional neural networks: 2D convolution

1	0	1
0	1	0
1	0	1

Filter / Kernel

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved Feature

Convolutional neural networks: multi-channel inputs

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



308

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-498

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



164

$$+ 1 = -25$$

$$\begin{array}{c} \uparrow \\ \text{Bias} = 1 \end{array}$$

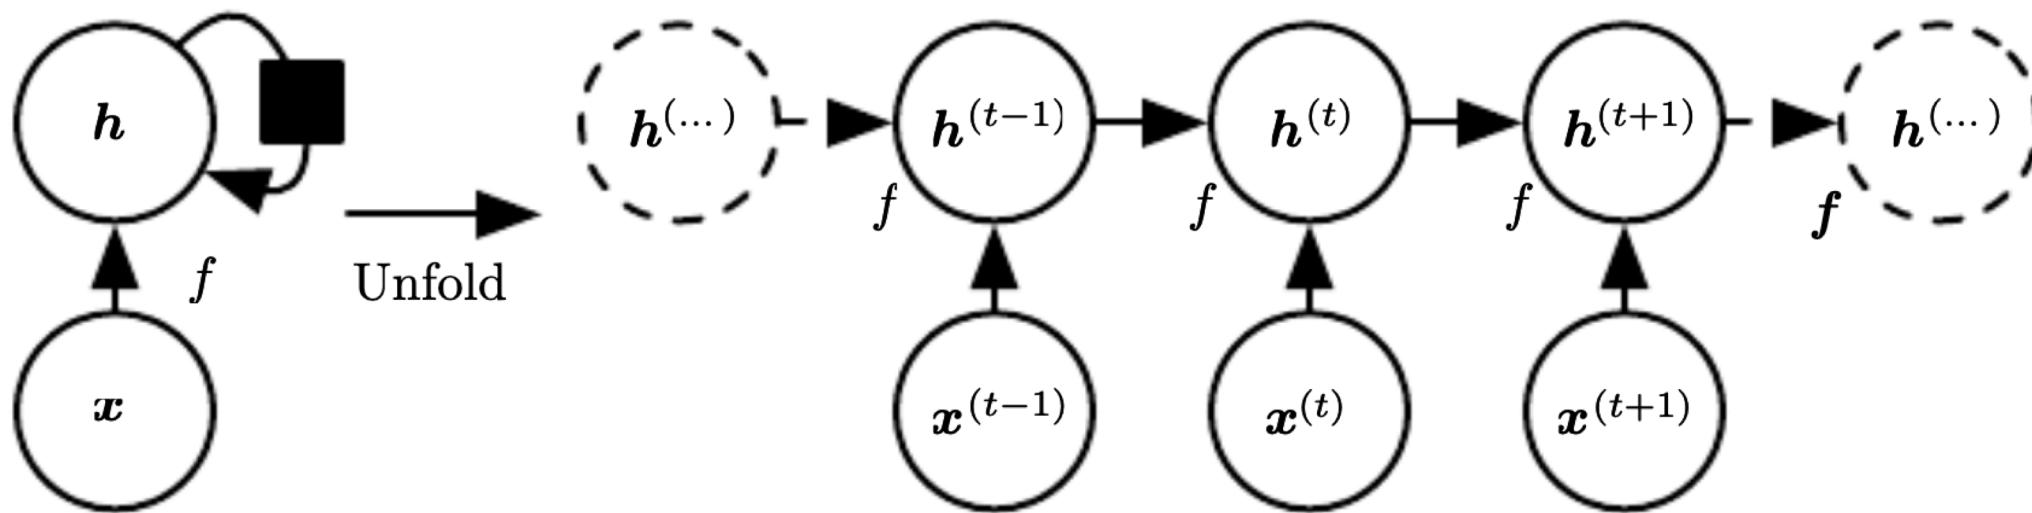
-25				...
				...
				...
				...
...

Output

Recurrent neural networks: state machine

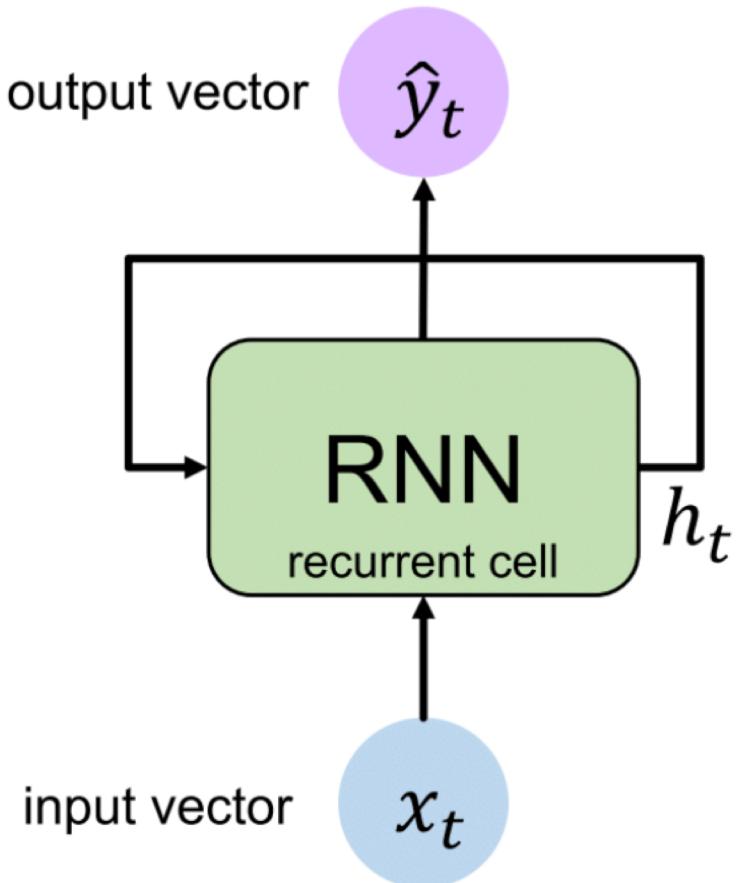
Recurrence relation w/ external input

$$\mathbf{s}^{(t)} = f(\mathbf{s}^{(t-1)}, \mathbf{x}^{(t)}; \boldsymbol{\theta})$$



Recurrent neural networks: back propagation through time

Idea: Learn sequential / temporal relationships



Apply a **recurrence relation** at every time step to process a sequence:

$$h_t = f_W(h_{t-1}, x_t)$$

new state function old state input vector at
parameterized by W time step t

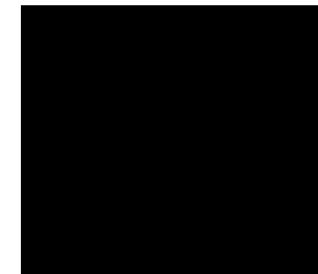
Note: the same function and set of parameters are used at every time step

Model Interpretability

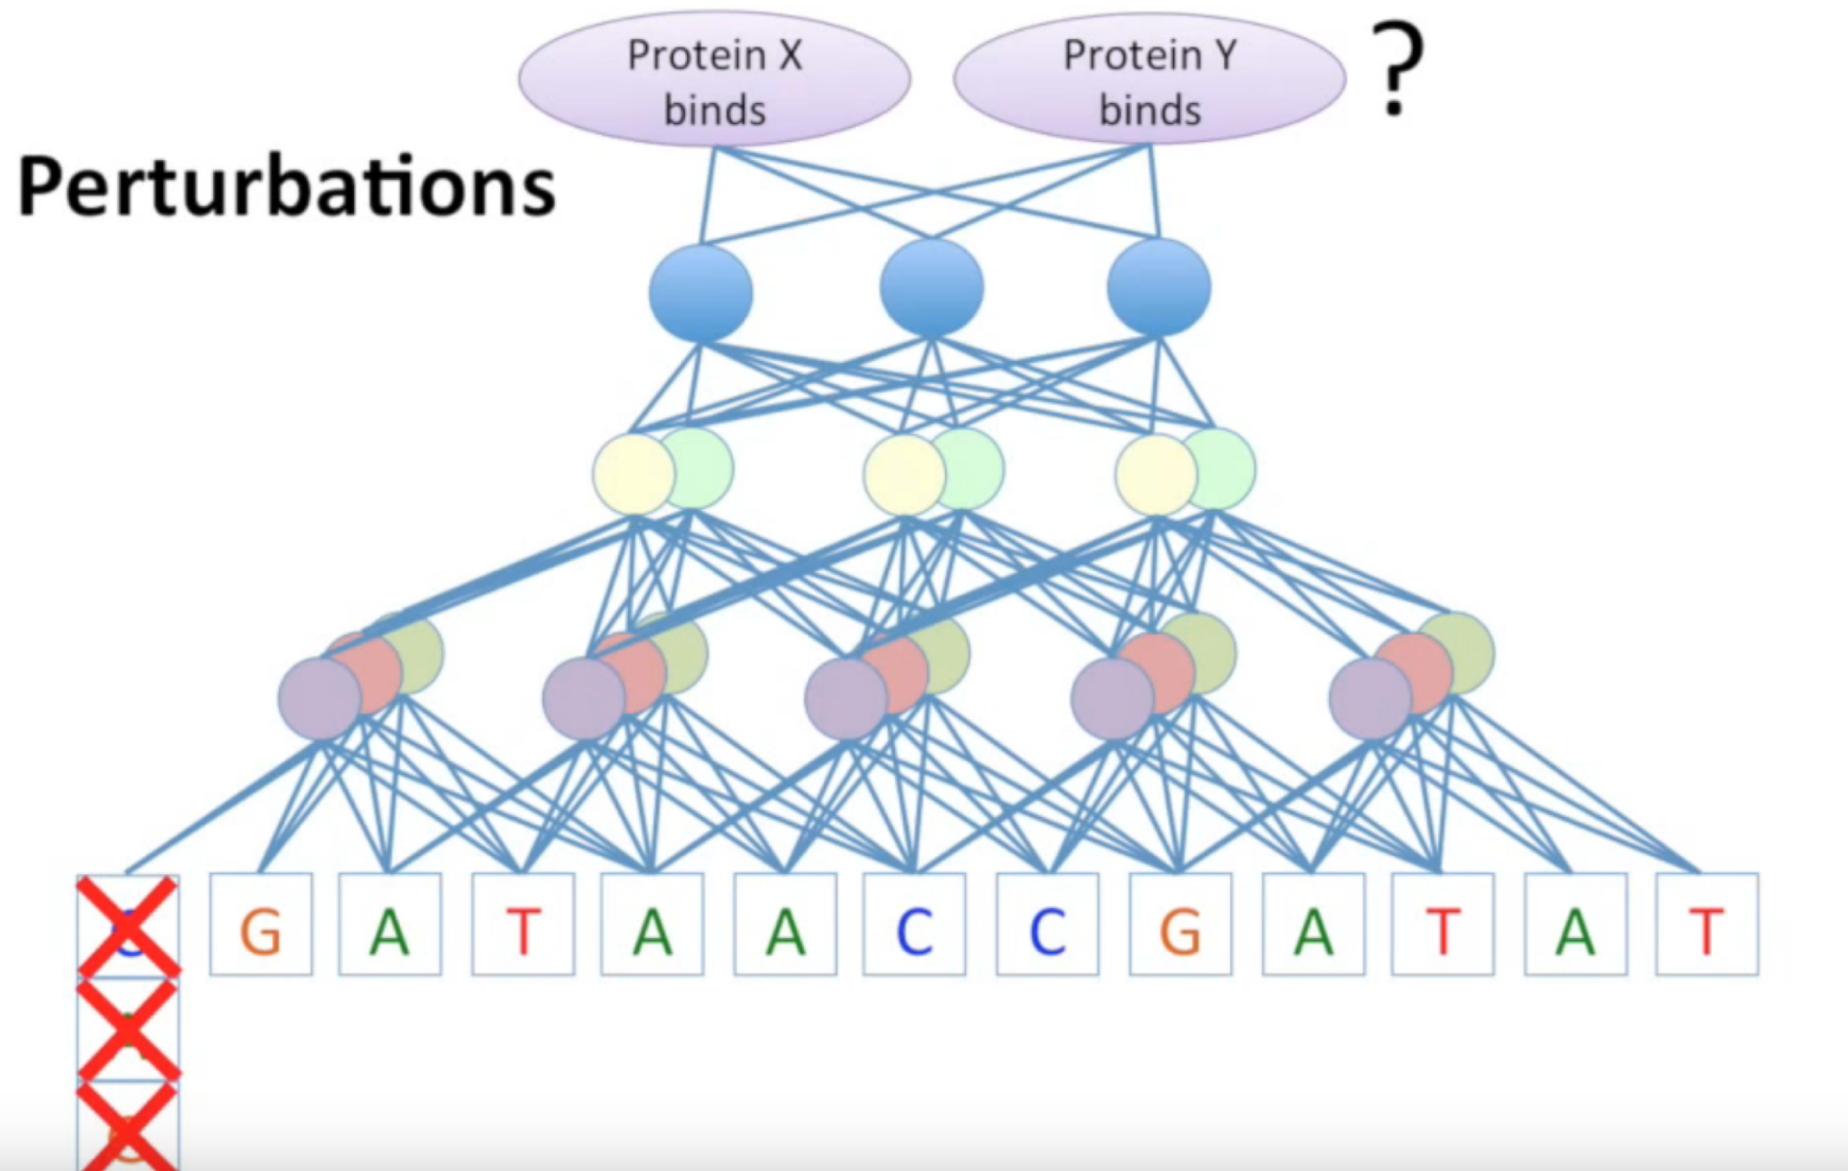
Slide credit to – Dylan Cable + Previous 6.874 Slides

Model interpretability: model interpretation overview

- Adoption of deep learning has led to:
 - Large increase in predictive capabilities
 - Complex and poorly-understood black-box models
- Imperative that certain model decisions can be interpretably rationalized
 - Ex: loan-application screening, recidivism prediction, medical diagnoses, autonomous vehicles
- Explain model failures and improve architectures
- Interpretability is also crucial in scientific applications, where goal is to identify general underlying principles from accurate predictive models



Model interpretability: black-box models



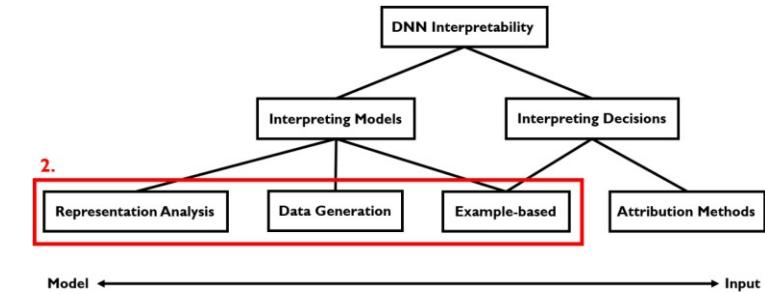
Model Interpretability: Four Basic Strategies

- Visualize the filters
- Measure the Gradient of node with respect to the input (Saliency Map)
- Choose an input that maximizes a node
- Perturb Input and See the effect on output
- Others!

Model Interpretability: Four Basic Strategies

- **Visualize the filters**
- Measure the Gradient of node with respect to the input (Saliency Map)
- Choose an input that maximizes a node
- Perturb Input and See the effect on output
- Others!

Types of DNN Interpretability



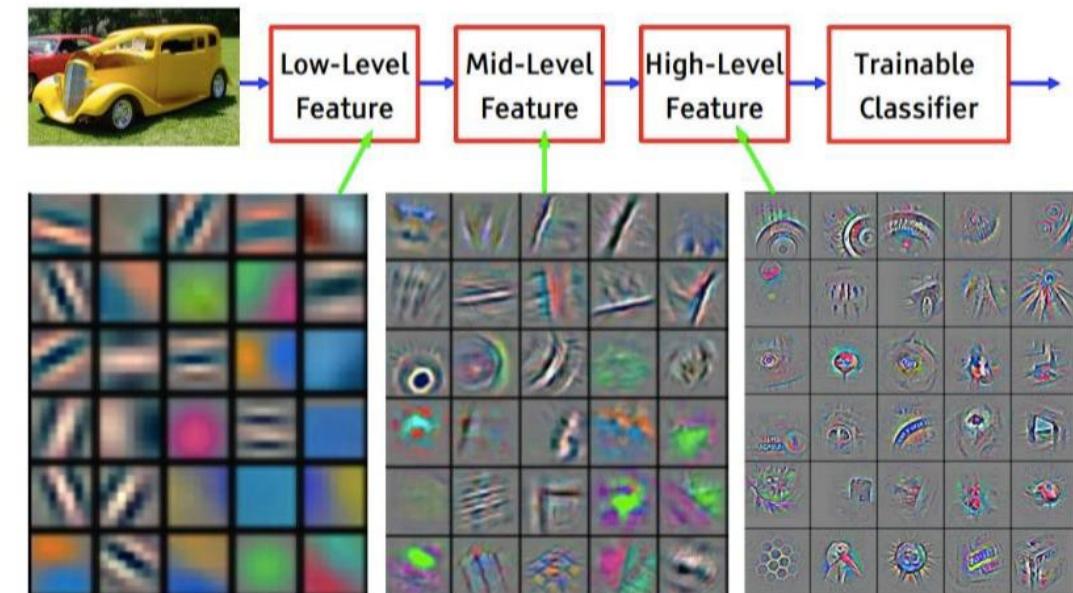
Weight Visualization

- Filter visualization in Convolutional Neural Networks
- Can understand what kind of features CNN has learned
- Still too many filters!

Surrogate Model

Data Generation

Example-based



Model interpretability: saliency maps

Visualizing saliency

The procedure is related to back-propagation, but in this case the optimization is performed with respect to the input image, while the weights are fixed to those found during the training stage.

More concretely,

$$S_c(I) \approx W^T I + b$$

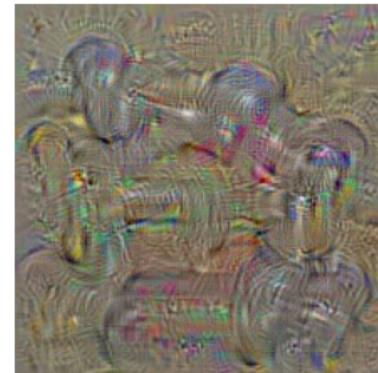
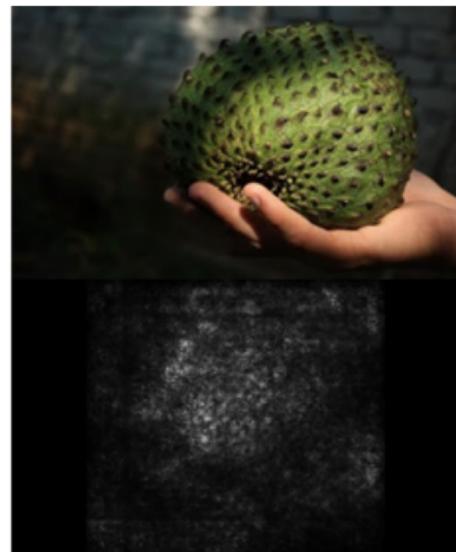
Where,

$$W = \frac{\partial S_c}{\partial I} \Big|_{I_i}$$

Optimization by gradient ascent

More formally, let $S_c(I)$ be the score of the class c , computed by the classification layer of the ConvNet for an image I . We would like to find (via backprop) an L2-regularised image, such that the score S_c is high,

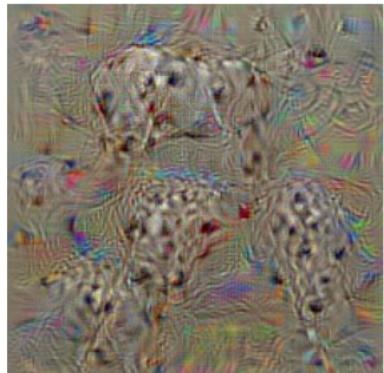
$$\arg \max_I S_c(I) - \lambda \|I\|_2^2$$



dumbbell

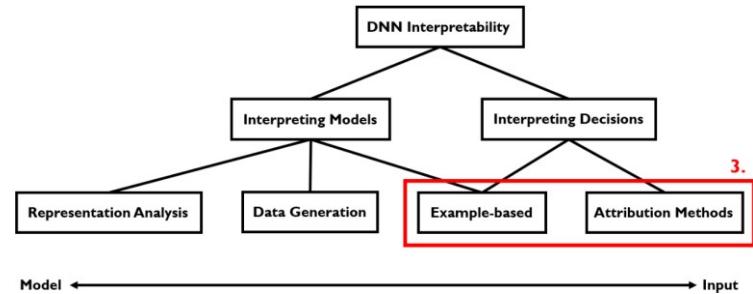


cup



dalmatian

Types of DNN Interpretability



Example-based

Attribution Methods

Gradient Based

Backprop. Based

- Which training instance influenced the decision most?
- Still does not **specifically highlight** which features were important
- Influence functions for interpreting black-box methods. Fragility of NN model interpretation.

'Sunflower': 59.2% conf.



Original

Influence: 0.09



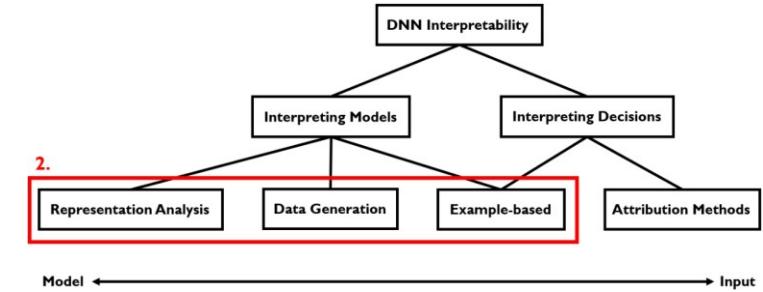
Influence: 0.14



Influence: 0.42



Types of DNN Interpretability



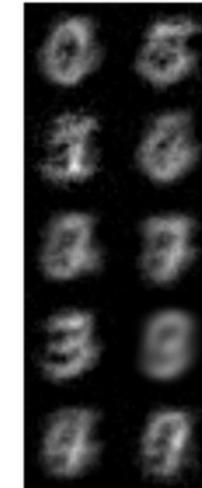
Weight Visualization

Surrogate Model

Data Generation

Example-based

simple AM
(initialized
to mean)

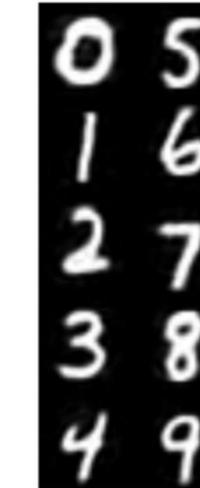


*Original
Activation
Maximization*

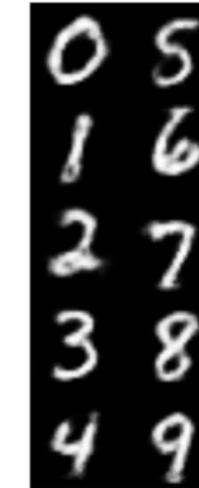
simple AM
(init. to
class
means)



AM-density
(init. to
class
means)



AM-gen
(init. to
class
means)



*Constrained
Activation
Maximization*

Observation: Connecting to the **data** leads to **sharper** visualizations.